

---

# Introduction to Computer Vision

## Lecture 1. Course Introduction

31.10.22

Mikhail Belyaev

# Outline

- **Info about the course**
- Computer Vision problems & applications
- Convolutions
- *Medical Computer Vision*

# Info about the course

---

## Term 2, 3 credit course.

- Blended mode.
- Mon: 12:30-15:30, Fri:9:00-12:00
  - Lectures: the first ~90 minutes (with recording).
  - Practical exercises to work at home (provided after the lecture).
  - Online review of the exercises with QA after the next lecture: the remaining 60-90 minutes.
- 12 lectures + hands-ons
- Final project

## Grading

- 3x20% - homeworks;
- 40% - final project.

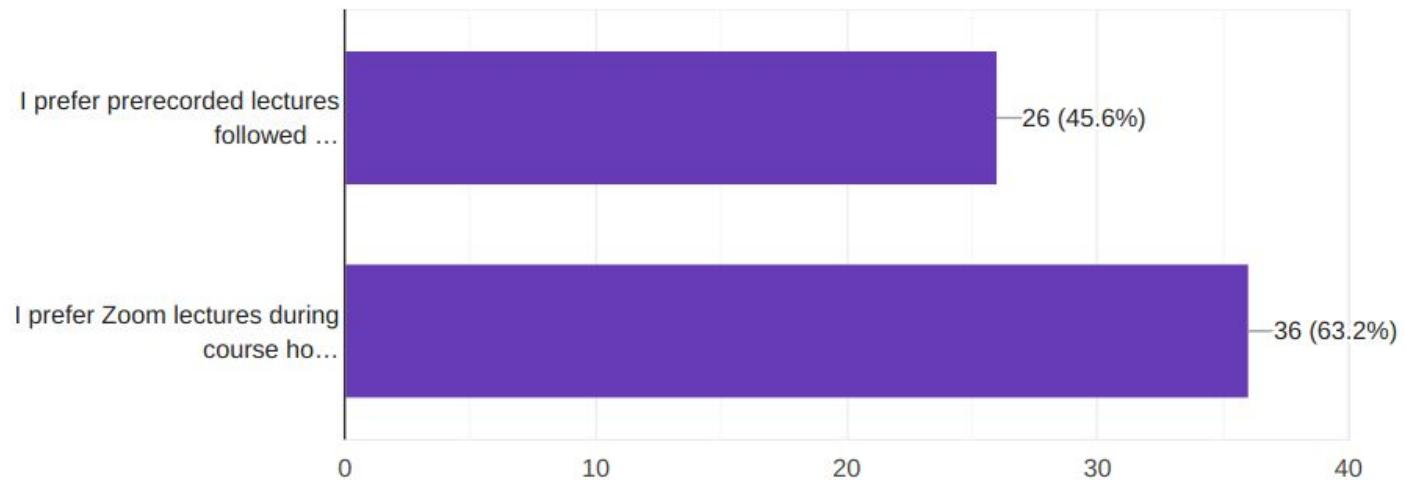
# Info about the course

---

## An old questionary about the format

Lectures: how to spend official course hours

57 responses



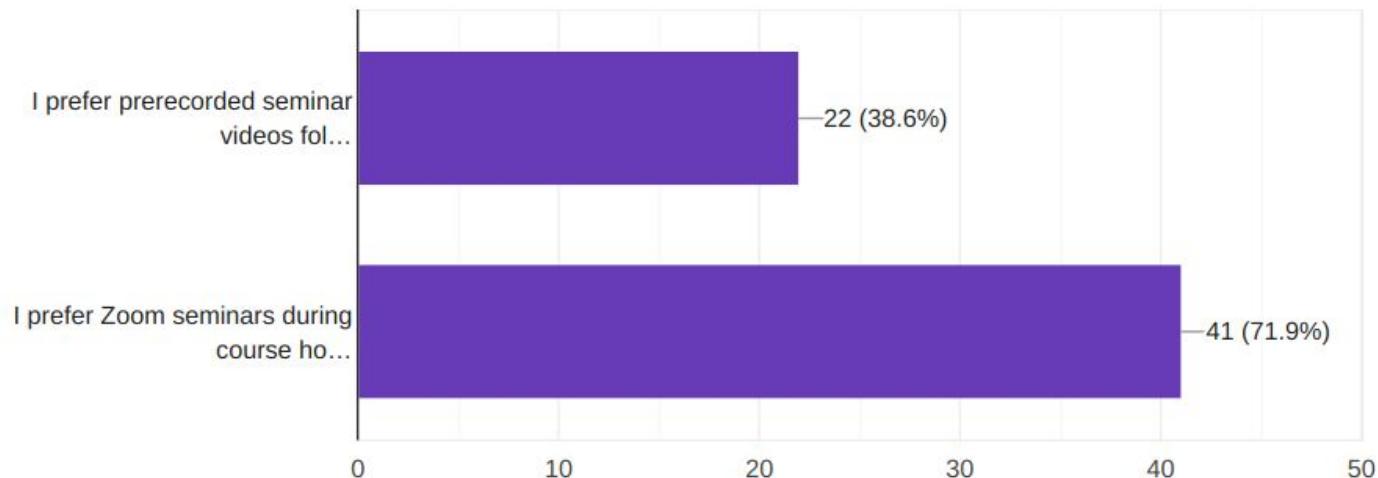
# Info about the course

---

## An old questionary about the format

Seminars: how to spend official course hours

57 responses



# Info about the courses

---

## Course outcomes

### Knowledge

1. Statements of all major computer vision problems.
2. Mathematical details of the most important computer vision algorithms.

### Skills

1. Select an appropriate method for solving particular computer vision problems.
2. Apply computer vision libraries.
3. Solve real-life computer vision problems.

### Knowledge

1. Implementing computer vision basic algorithms and complex pipelines

# Course team

---

**Mikhail  
Belyaev,  
Instructor**



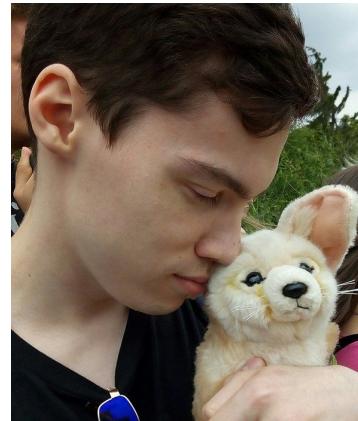
**Natalia  
Soboleva,  
TA**



**Anton  
Vasilyuk,  
TA**



**Boris  
Shirokikh,  
TA**



# Course structure

---

## Part 1. Classical CV algorithms. 9 lectures

1. Template matching. Edge detection.
2. Binary images. Morphological operations, contours in OpenCV.
3. Color spaces. Segmentation
4. Hough transform. Projective transform. RANSAC.
5. Object detection & classification: Haar cascade, Histogram of Oriented Gradients.
6. Keypoint detection. Blob detection. Scale invariance.
7. Image descriptors.
8. Optical flow
9. Object tracking.

## Part 2. Convolutional neural networks intro. 2 lectures.

1. CNN basics. CNN for classification, detection & segmentation.

# Course Evaluation - 2020

---

## Standard Course Evaluation - AY2020/21

1. Course objectives were clear to me	Strongly Agree	49 (71.01%)
	Agree	18 (26.09%)
	Disagree	2 (2.90%)
2. Key concepts and theories were well explained by the Course Instructor(s)	Strongly Agree	46 (65.71%)
	Agree	24 (34.29%)
3. Course content was difficult enough to be challenging	Strongly Agree	37 (53.62%)
	Agree	31 (44.93%)
	Disagree	1 (1.45%)

# Course Evaluation - 2020

---

## Standard Course Evaluation - AY2020/21

1. Course objectives were clear to me	Strongly Agree	49 (71.01%)
	Agree	18 (26.09%)
	Disagree	2 (2.90%)
2. Key concepts and theories were well explained by the Course Instructor(s)	Strongly Agree	46 (65.71%)
	Agree	24 (34.29%)
3. Course content was difficult enough to be challenging	Strongly Agree	37 (53.62%)
	Agree	31 (44.93%)
	Disagree	1 (1.45%)

Everything was cool but why do we need the course?

# Course Evaluation - 2020

---

## Standard Course Evaluation - AY2020/21

1. Course objectives were clear to me	Strongly Agree	49 (71.01%)
	Agree	18 (26.09%)
	Disagree	2 (2.90%)

2. Key co	.71%)
	.29%)

3. Course	.62%)
	.93%)
	5%)

## DISCLAIMER

This course is not about Deep learning!

Everything was cool but why do we need the course?

# Course Evaluation - 2020

---

## Why do we need classical computer vision methods in 2021?

1. Many ideas in modern deep learning methods are based on classical computer vision methods.
2. Deep learning methods are data hungry. Using classical CV method you can create a prototype faster. Also, some unsupervised approaches can be used as to generate annotation candidates

# Plagiarism policy

---

Rules are simple:

1. If you use external code, please add the link
2. Homeworks are individual assignments. If one homework is submitted twice, both student got  $\frac{1}{2}$  of the score.
3. Final projects are team-based. However, do not submit a random project from internet! It will hurt your overall score!

# Homework 1

Do you know this game? If yes, what is the most boring part?



# Homework 1

---

And counting scores at the end of the game is really boring!

The goal of HW1 is to create an automated scoring system.



# Homework 1

---

The goal of HW1 is to create an automated scoring system for Ticket to ride!

## Challenges

1. Detect all town markers and trains
2. Split cells vs trains
3. Estimate the colors (e.g. take into account different light conditions)



# Homework 1 - grading

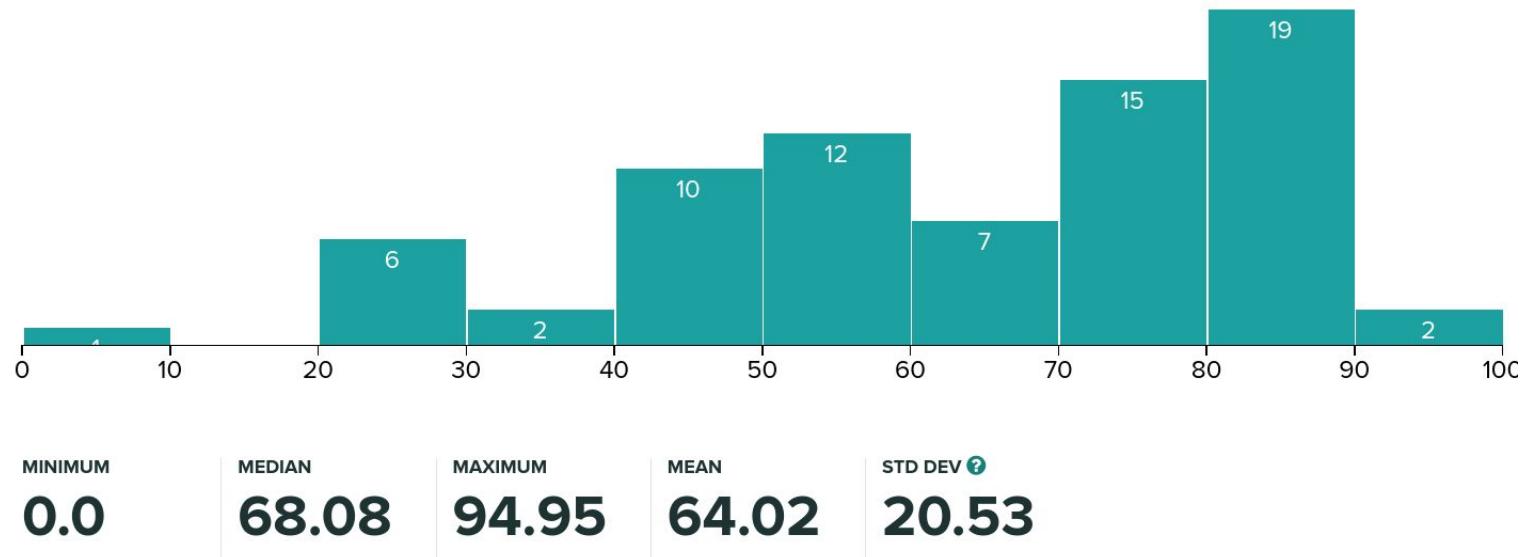
We will try to use gradescope as usual.

- 2.5 points for city detection
- 5.0 points per image for detecting trains.
- 2.5 points for the final game score.

10 points for each of the 10 test images will result in 100 points (max).

Review Grades for **Homework 1: Ticket to Ride**

● GRADES PUBLISHED



# Outline

- Info about the course
- **Computer Vision problems & applications**
- Convolutions
- *Medical Computer Vision*

# Computer Vision or Image processing?

---

What is the difference between Image Processing and Computer Vision? There are various definitions, but we'll stick with the following.

Computer Vision

- Image or Video -> Interpretation (understanding)

Image Processing:

- Image / Images -> A new Image / New Images

# CV problems: Image Classification



Image source: Krizhevsky et al. *ImageNet Classification with Deep Convolutional Neural Networks*

# CV problems: Object detection

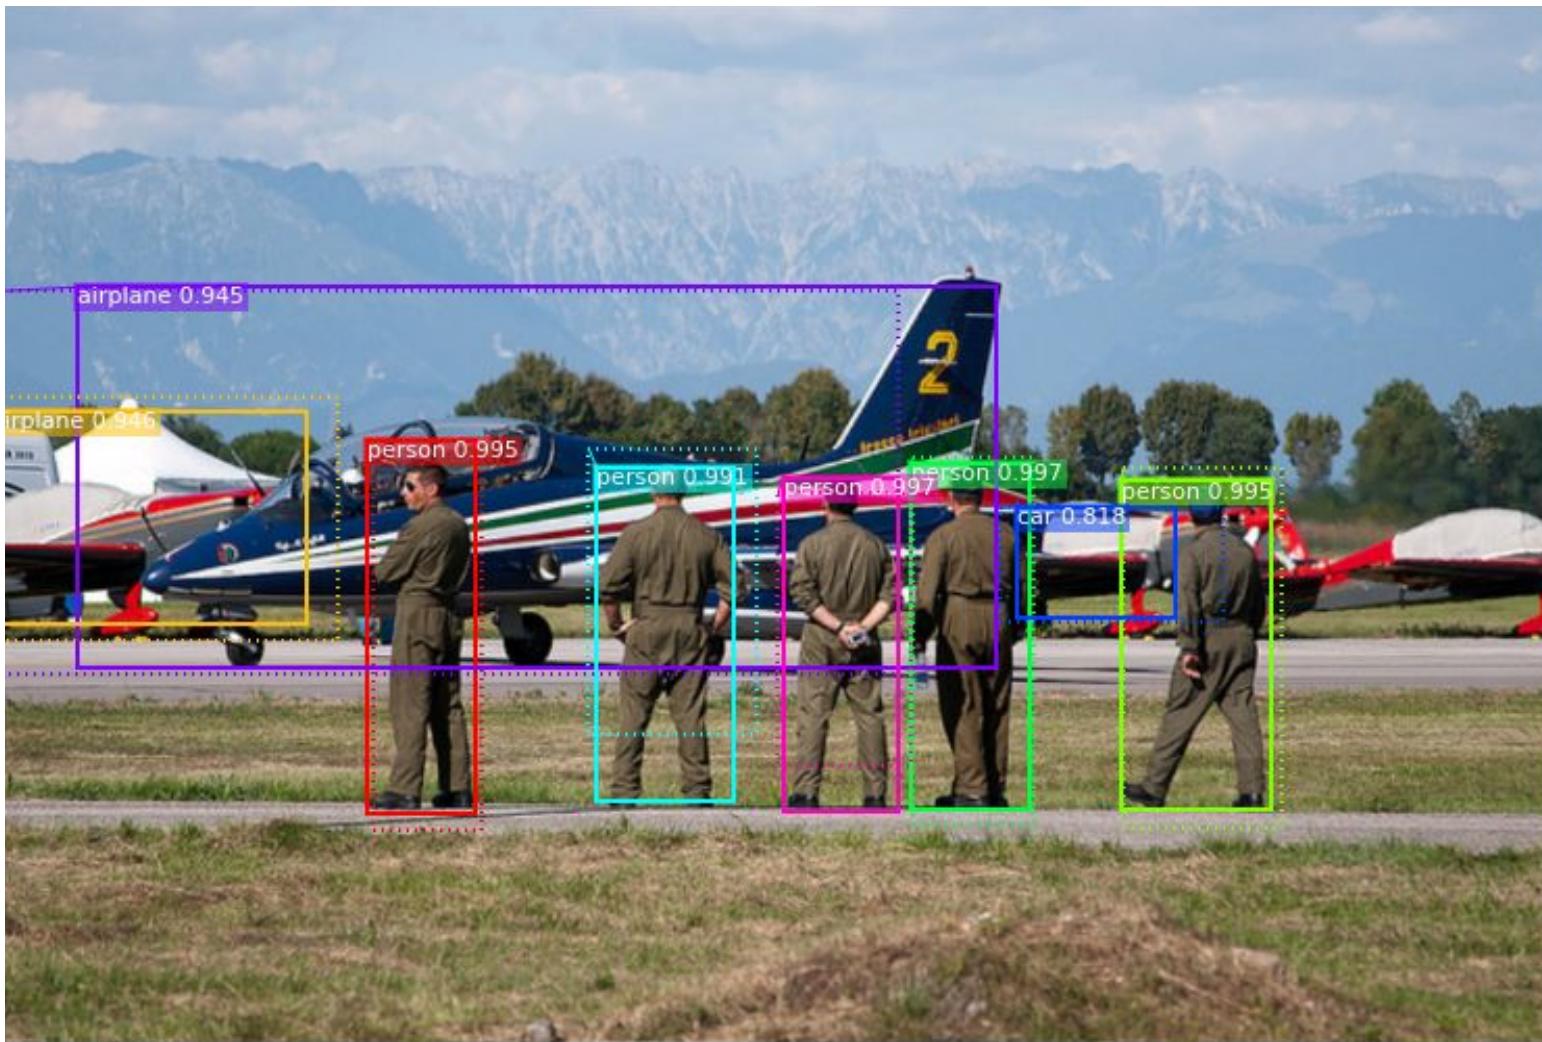
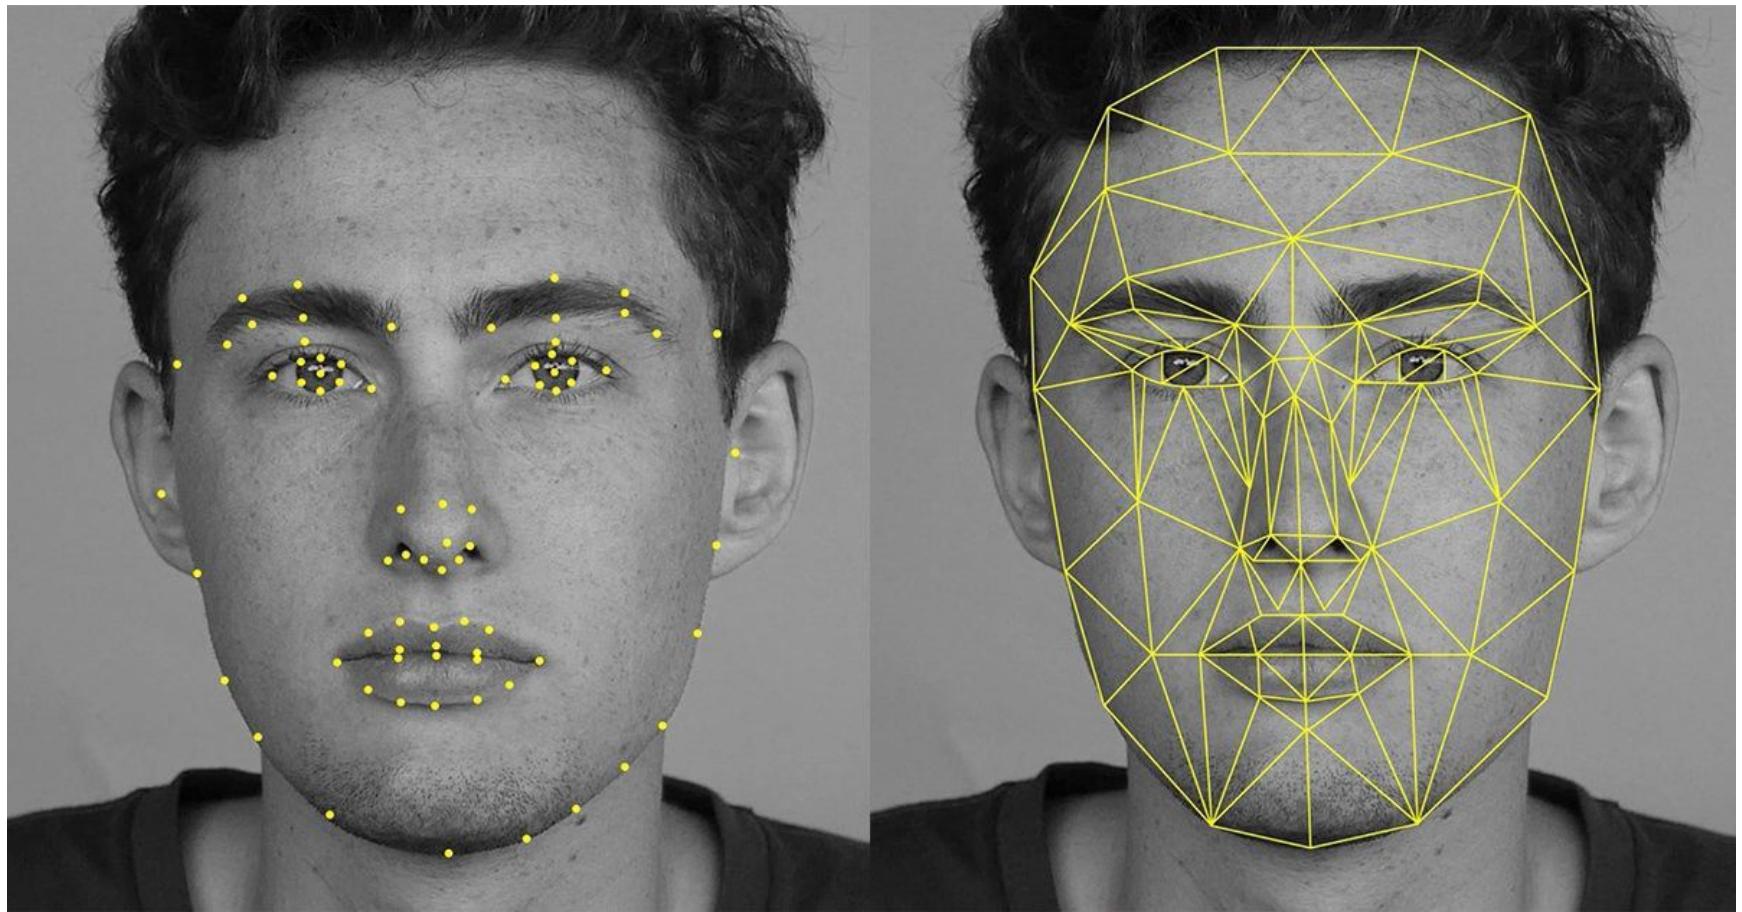


Image source: [https://github.com/matterport/Mask\\_RCNN](https://github.com/matterport/Mask_RCNN)

# CV problems: Keypoints detection

---



# CV problems: Segmentation



Image source: <https://www2.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/resources.html>

# CV problems: Semantic segmentation

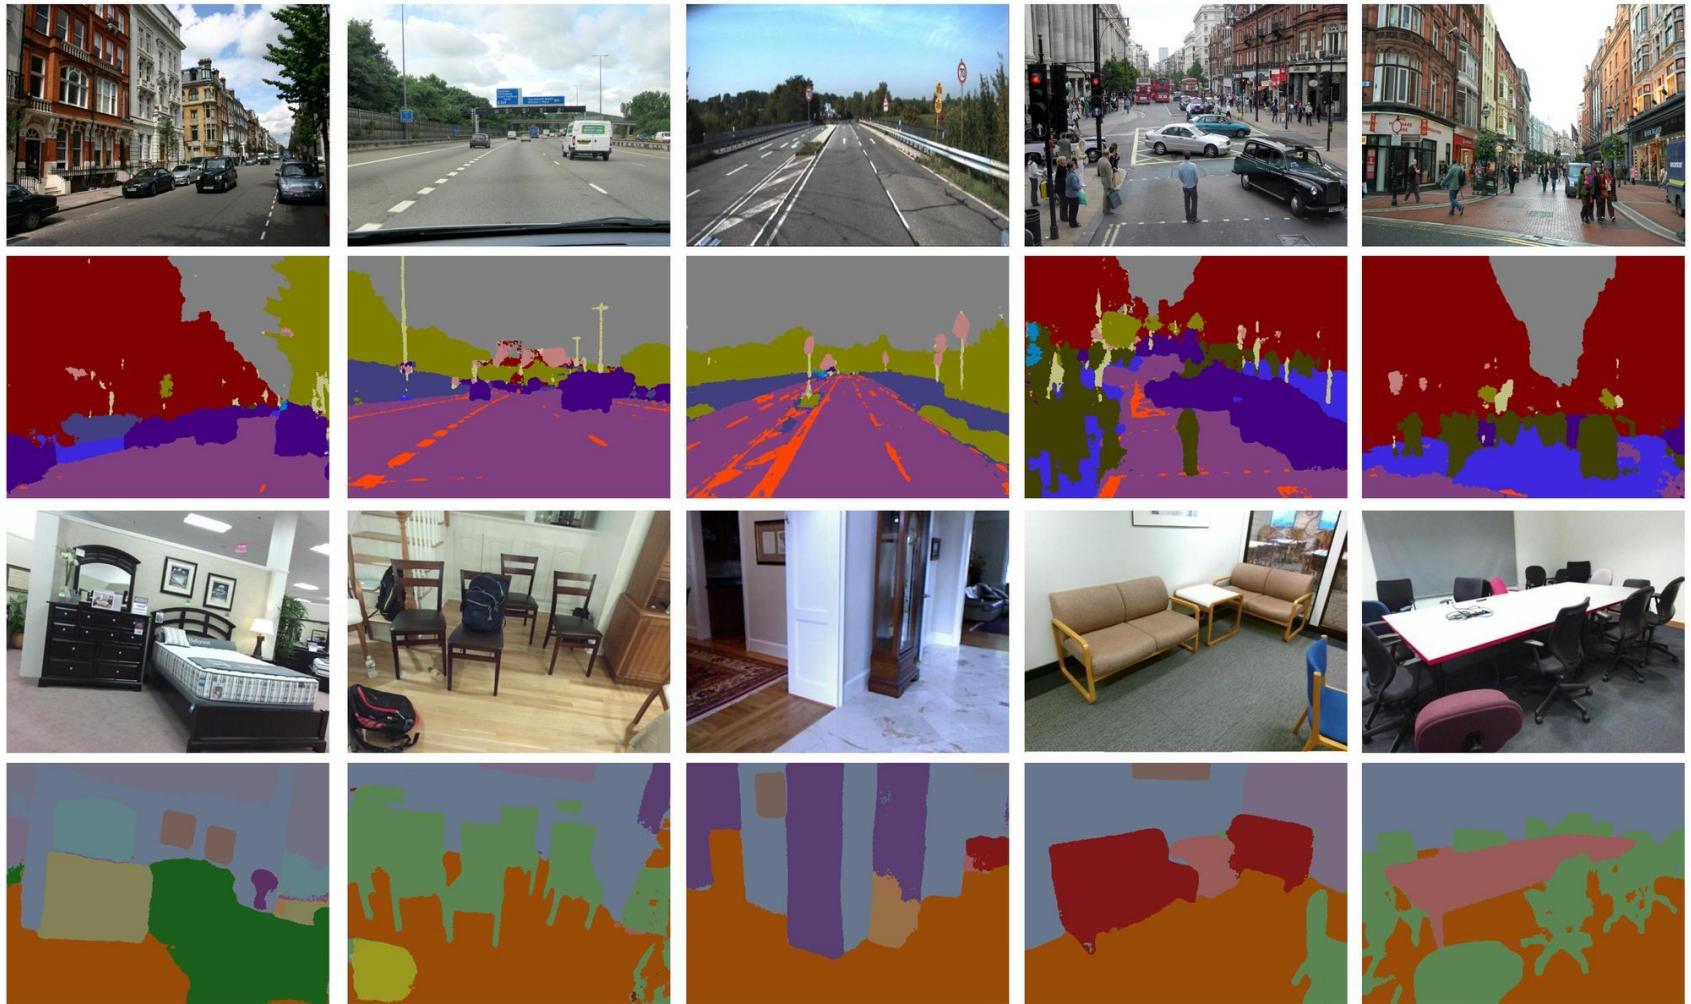


Image source: <https://mi.eng.cam.ac.uk/projects/segnet/>

# CV problems: Image Captioning

A person riding a motorcycle on a dirt road.



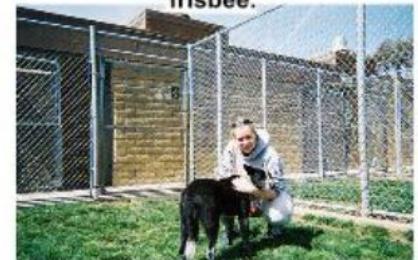
Two dogs play in the grass.



A skateboarder does a trick on a ramp.



A dog is jumping to catch a frisbee.



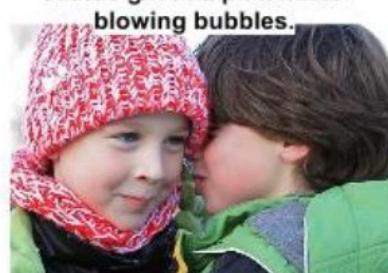
A group of young people playing a game of frisbee.



Two hockey players are fighting over the puck.



A little girl in a pink hat is blowing bubbles.



A refrigerator filled with lots of food and drinks.



A herd of elephants walking across a dry grass field.



A close up of a cat laying on a couch.



A red motorcycle parked on the side of the road.



A yellow school bus parked in a parking lot.



Image source: Vinyals et al. *Show and Tell: A Neural Image Caption Generator*

# CV problems: Depth Estimation

---

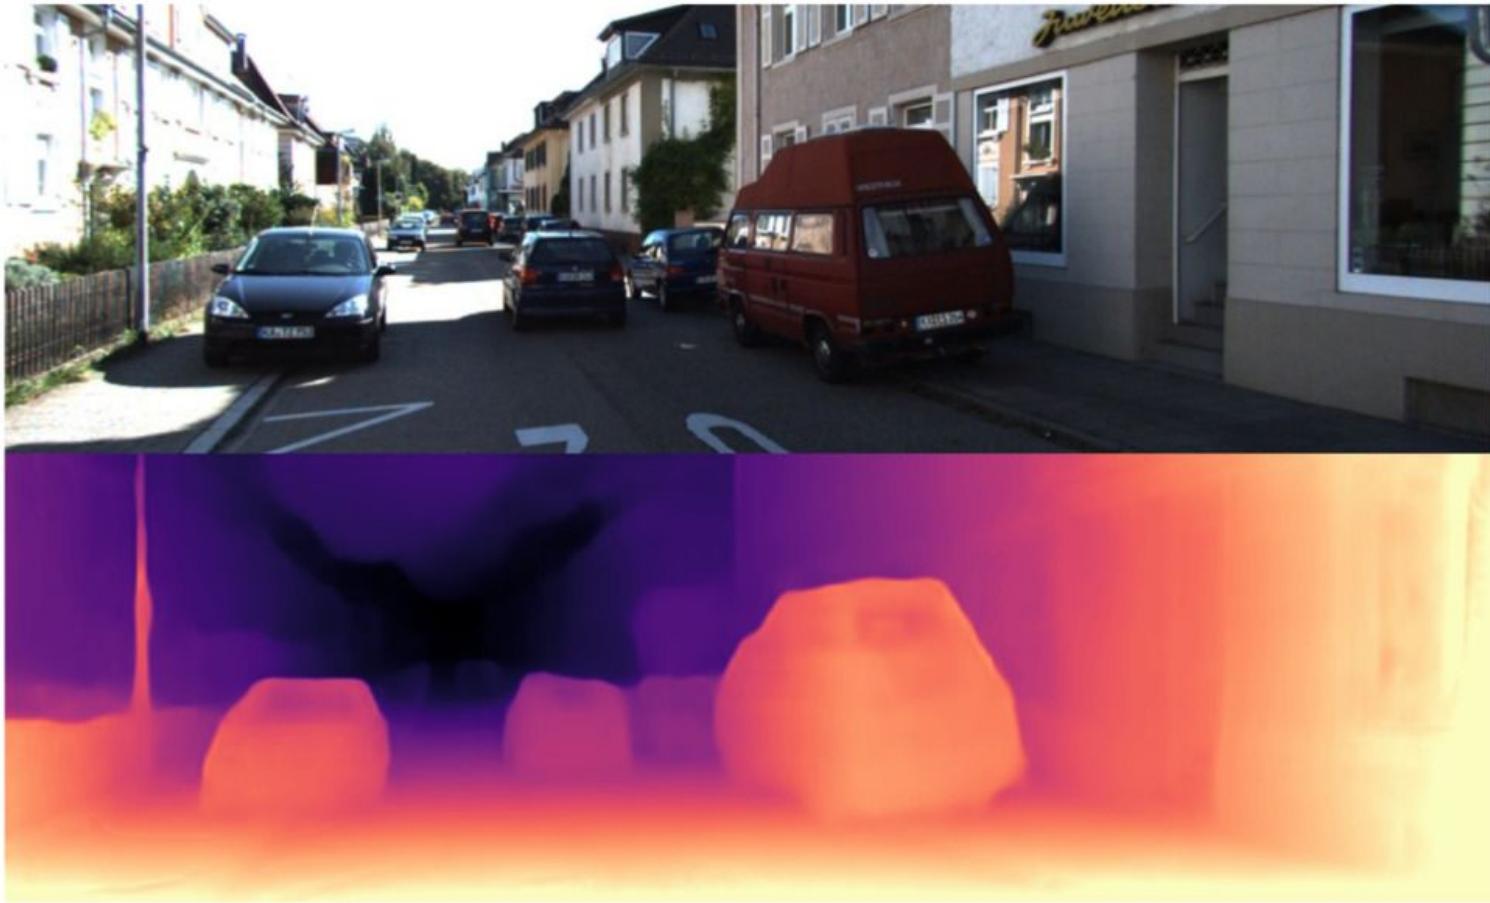
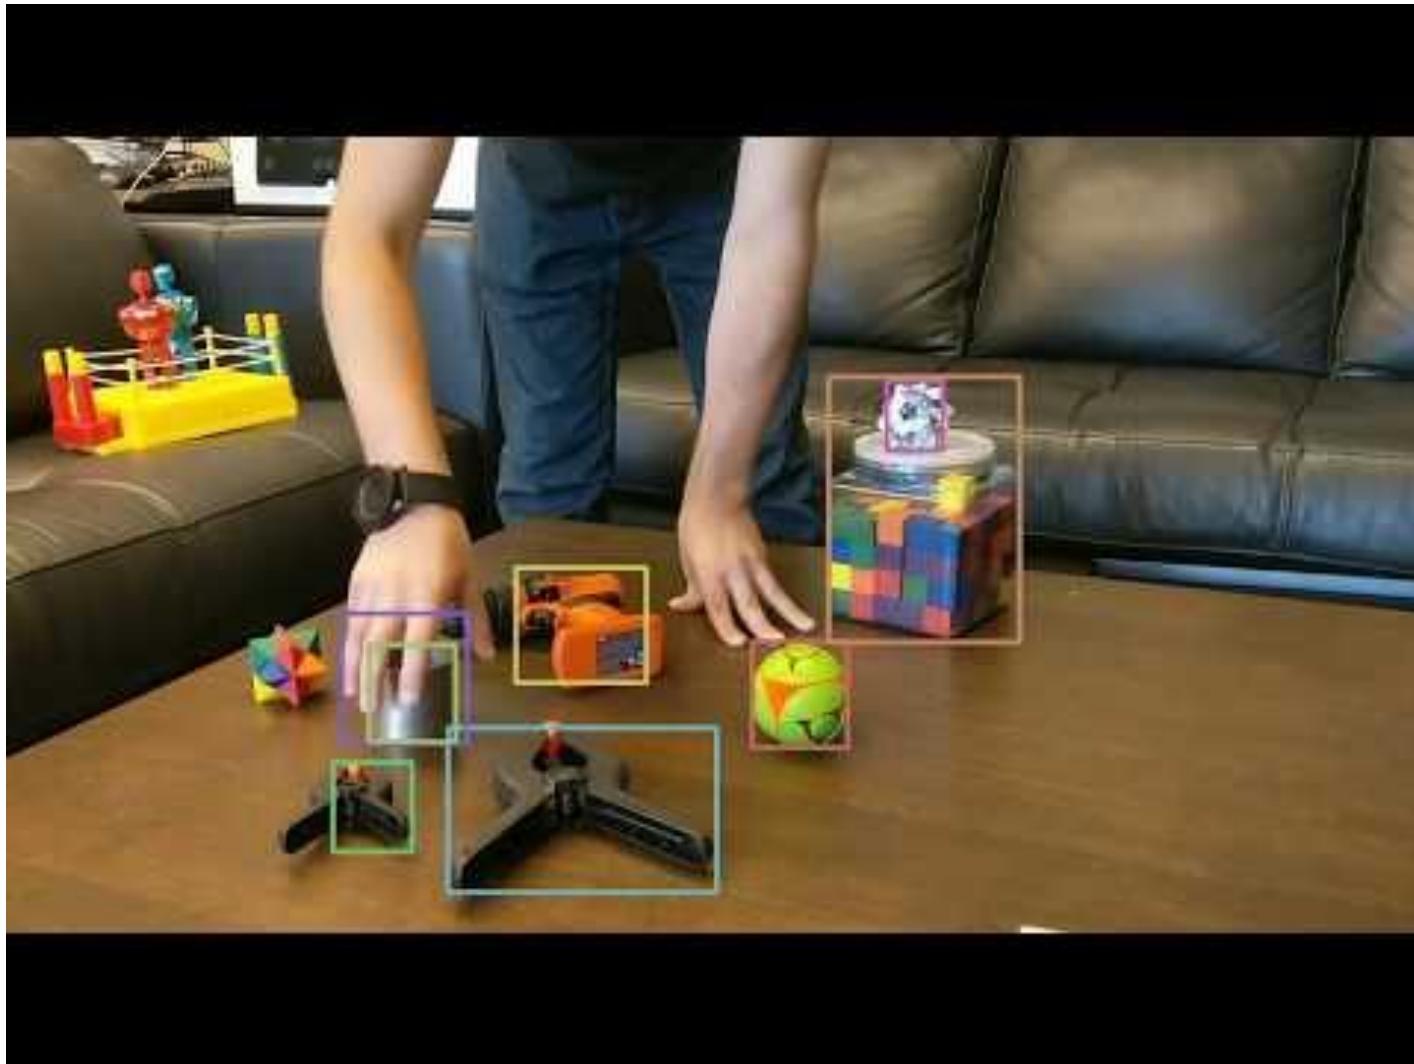


Image source: <https://github.com/nianticlabs/monodepth2>

# CV problems: Object tracking

---



Video source: <https://arxiv.org/abs/1705.06368>

# Image Processing problems: Super-resolution

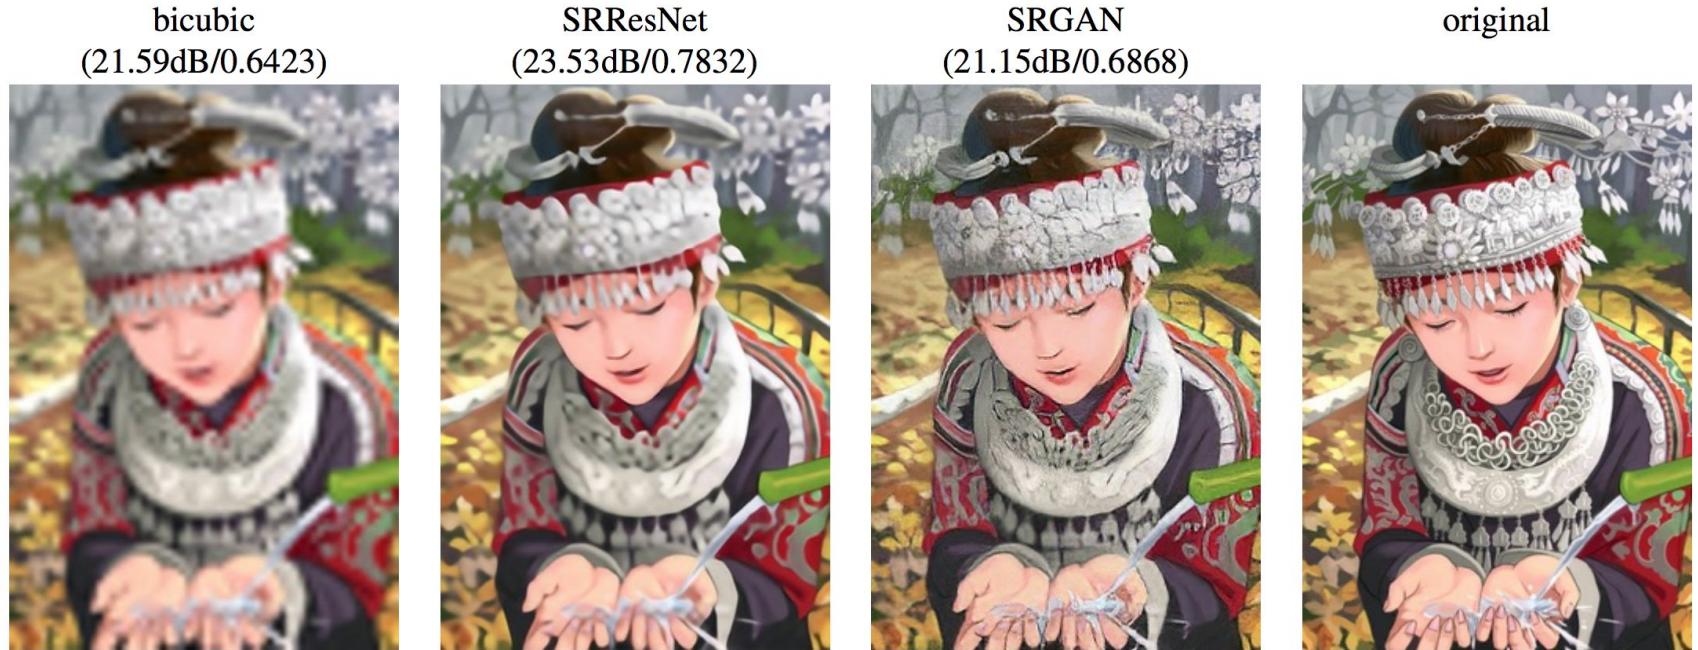


Figure 2: From left to right: bicubic interpolation, deep residual network optimized for MSE, deep residual generative adversarial network optimized for a loss more sensitive to human perception, original HR image. Corresponding PSNR and SSIM are shown in brackets. [4× upscaling]

Image source: Ledig et al. *Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network*

# Image Processing problems: Inpainting

---

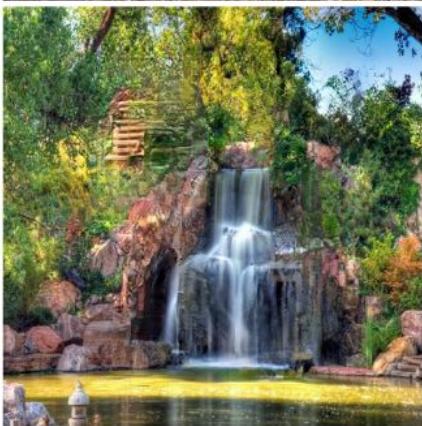
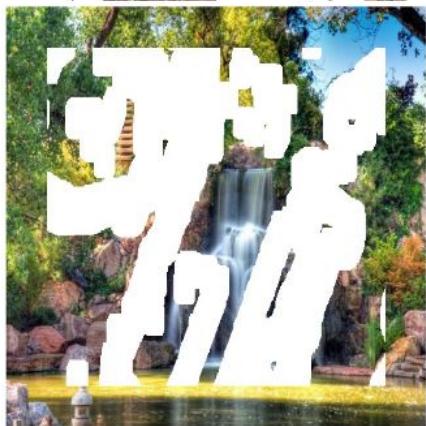
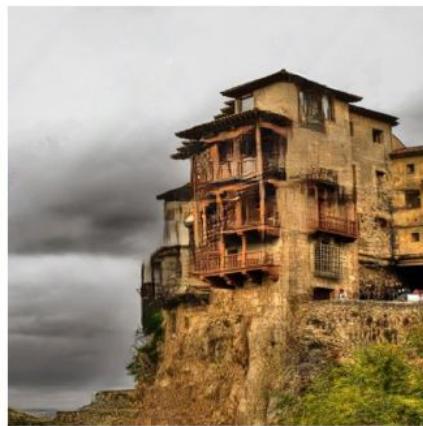


Image source: Liu et al. *Image Inpainting for Irregular Holes Using Partial Convolutions*

# Image Processing problems: Image generation

---



Image sources: Gravity, Jurassic Park, Avatar

# Image Processing problems: Style Transfer

Content image



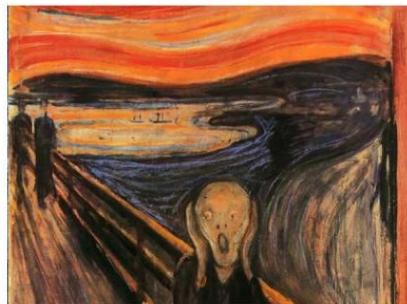
Style image



Output image



+



+



+



Image sources: <https://blog.godatadriven.com/images/how-to-style-transfer>

# Image Processing problems: Face generation

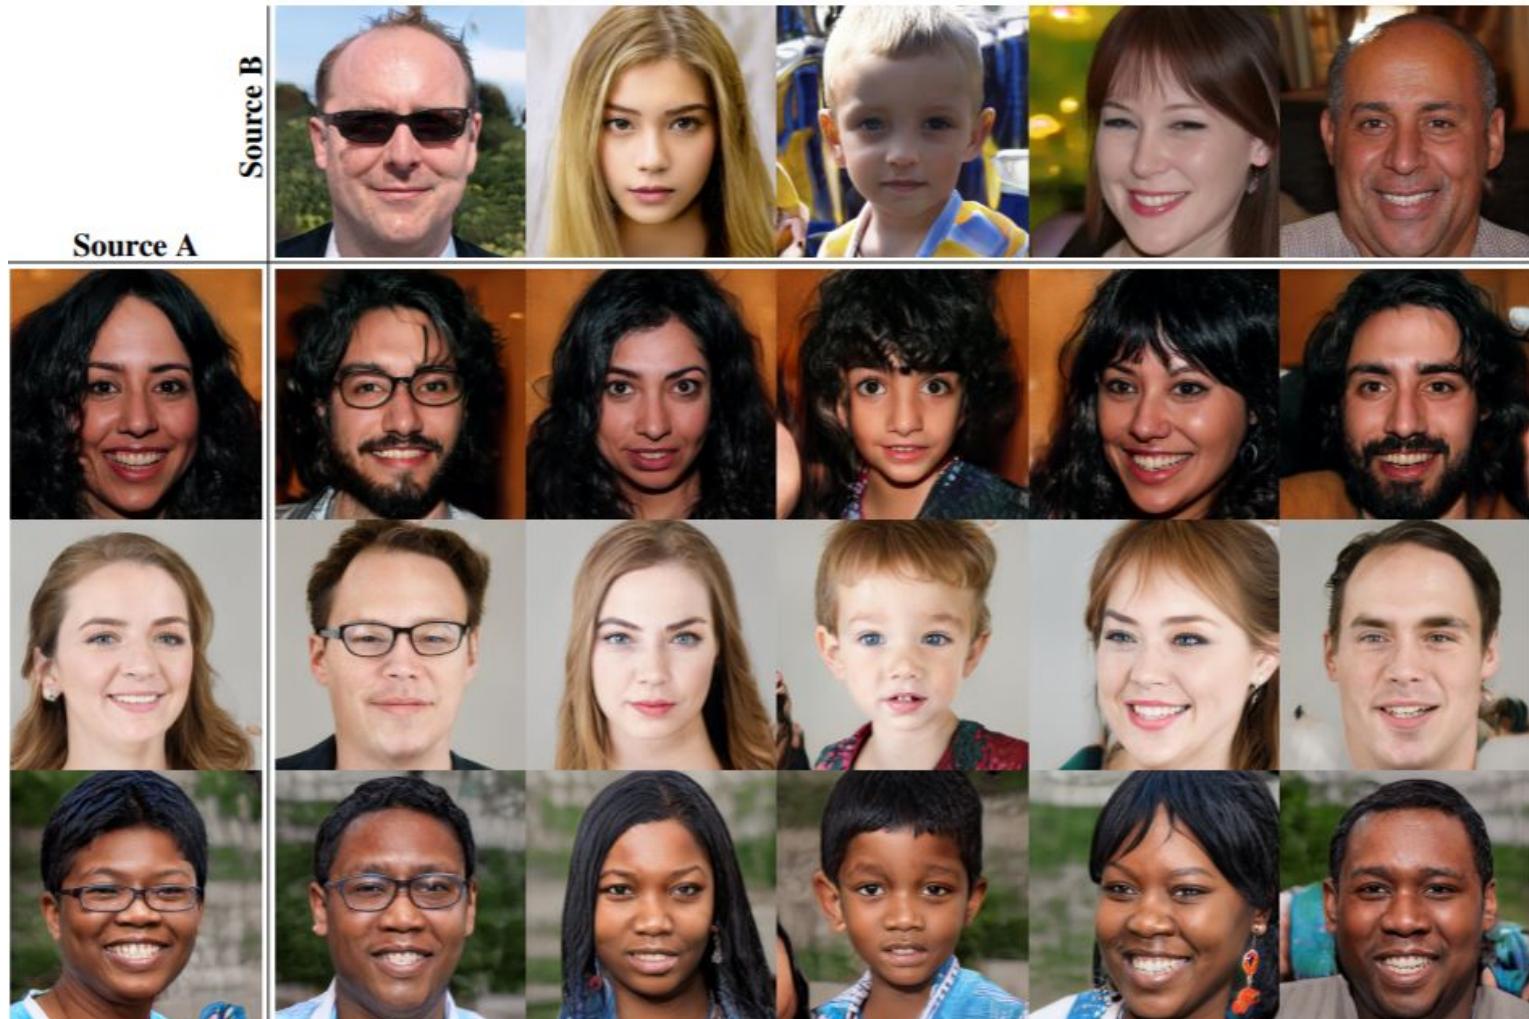


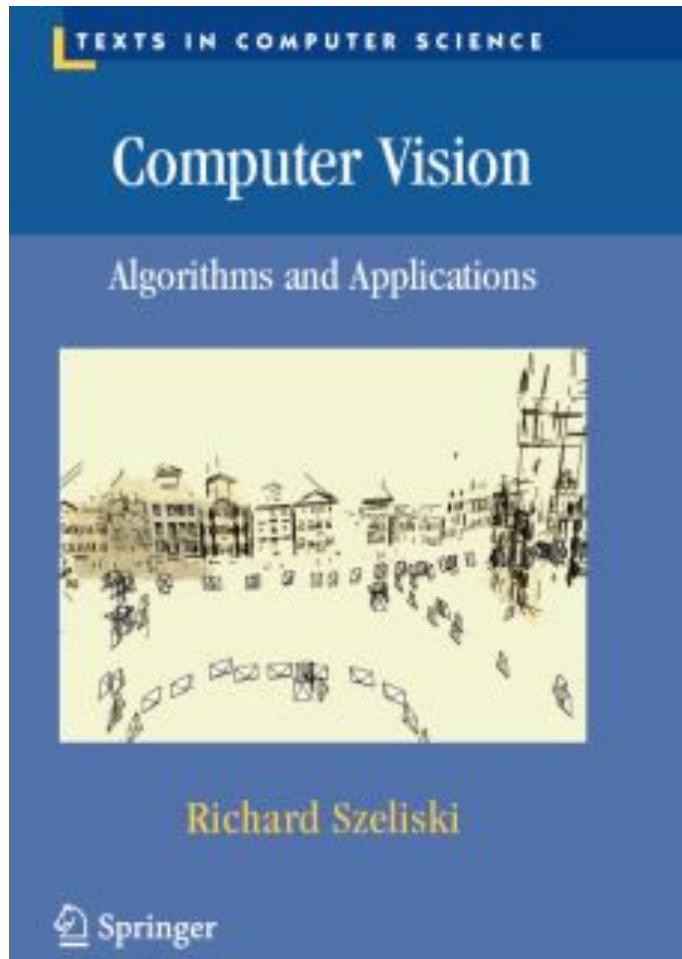
Image sources: <https://arxiv.org/pdf/1812.04948.pdf>

# Outline

- Info about the course
- Computer Vision problems & applications
- **Convolutions**
- *Medical Computer Vision*

# Reading

---



Good reading for this lecture.

The book is available [online](#)

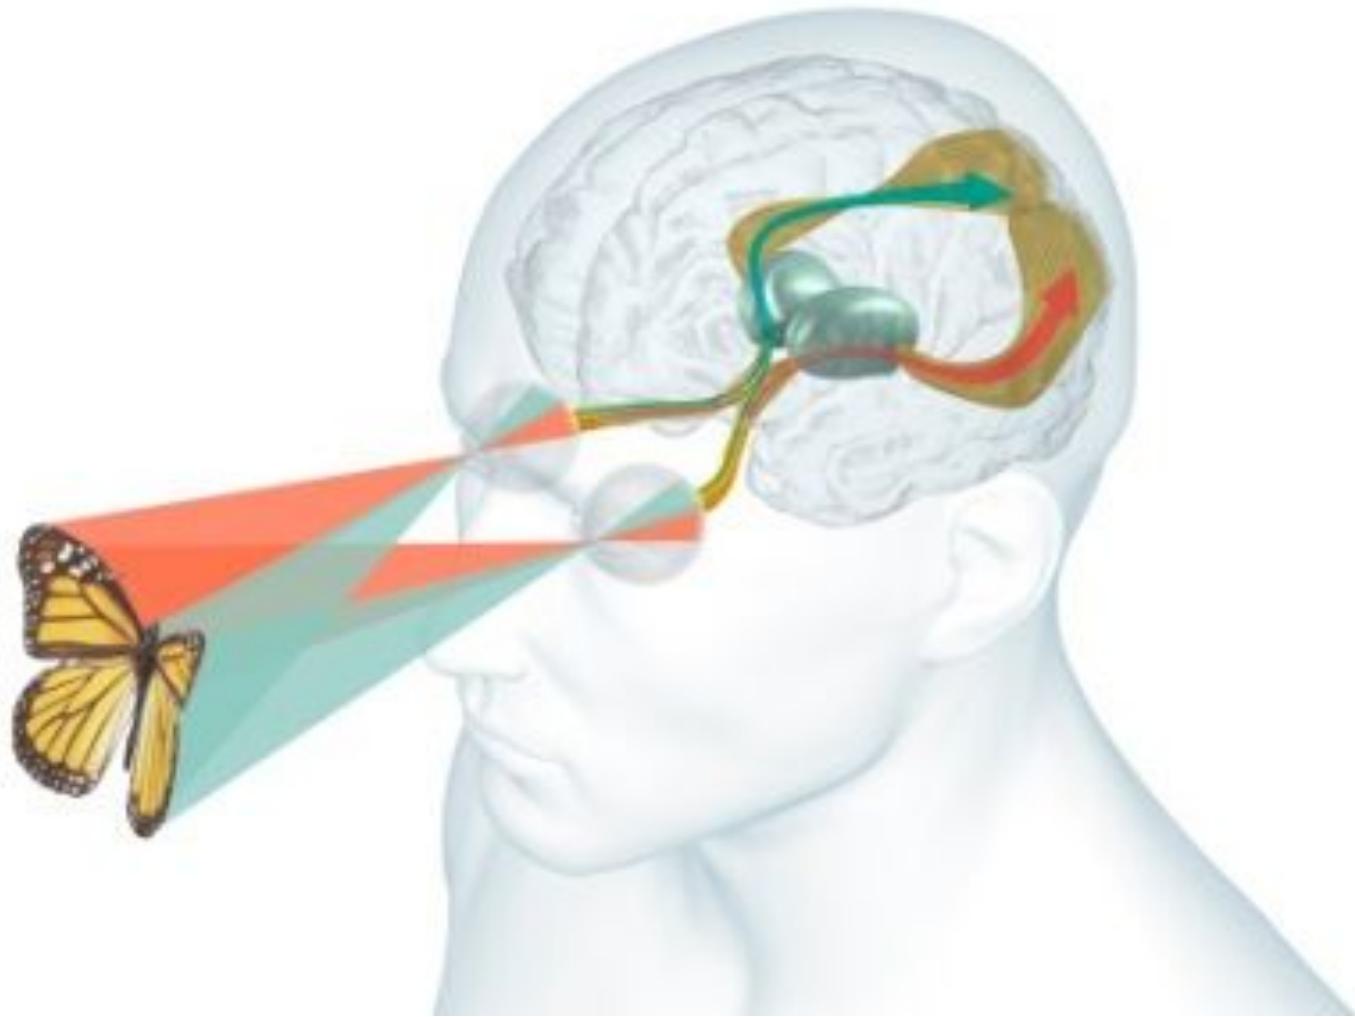
# What is an *image*?

---



# How do humans perceive images?

---



# How do humans perceive images?

---



What are the colors of this dress?

# How do humans perceive images?

---



Image perception is subjective: different people will see this dress (left) either as blue and black or as white and brown.

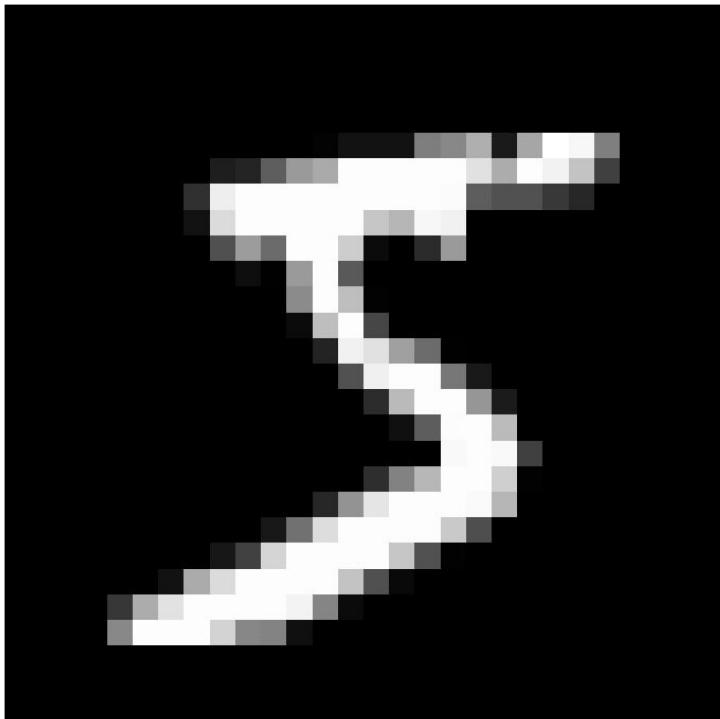
# How do humans perceive images?

---

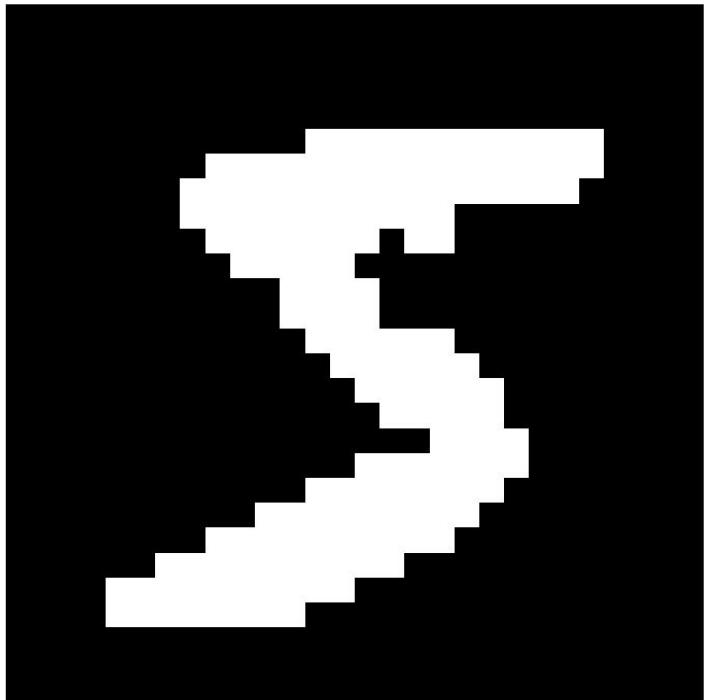


Image perception is subjective: do you see a young woman or an old one?

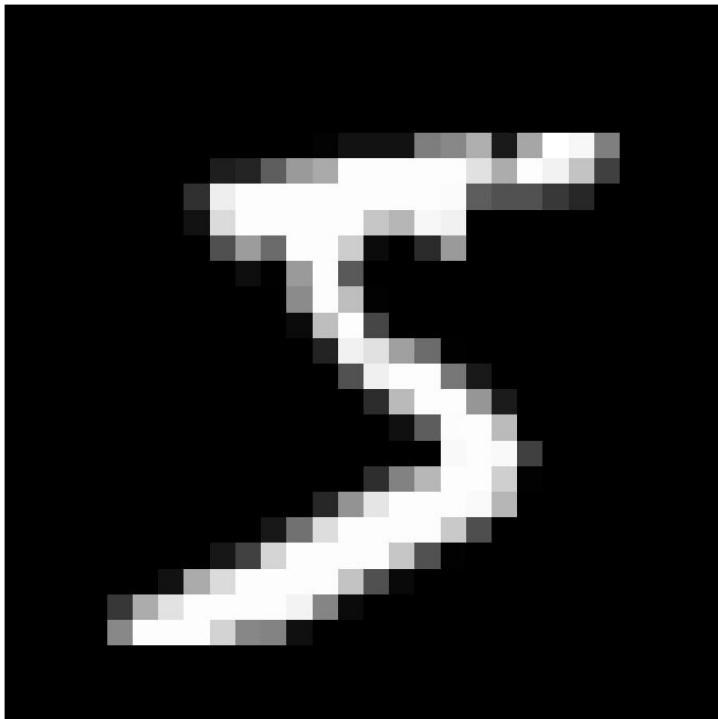
# How do machines store images?



# How do machines store images: Binary



# How do machines store images: Grayscale



Integer values from 0 to 255 (1 byte per pixel)

# How do machines store images: RGB

Red



Green



Blue

# How do machines store images: RGB

Red

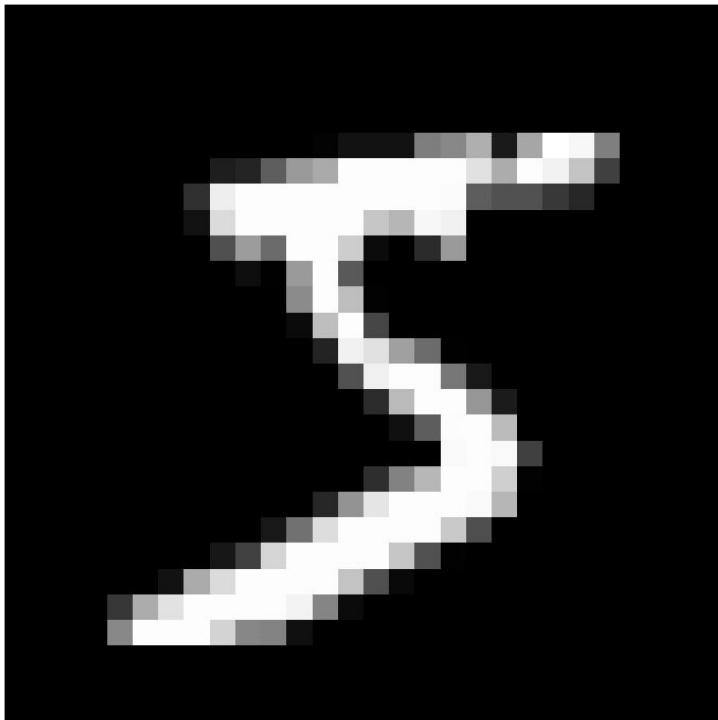


Green



Blue

# How do machines store images?

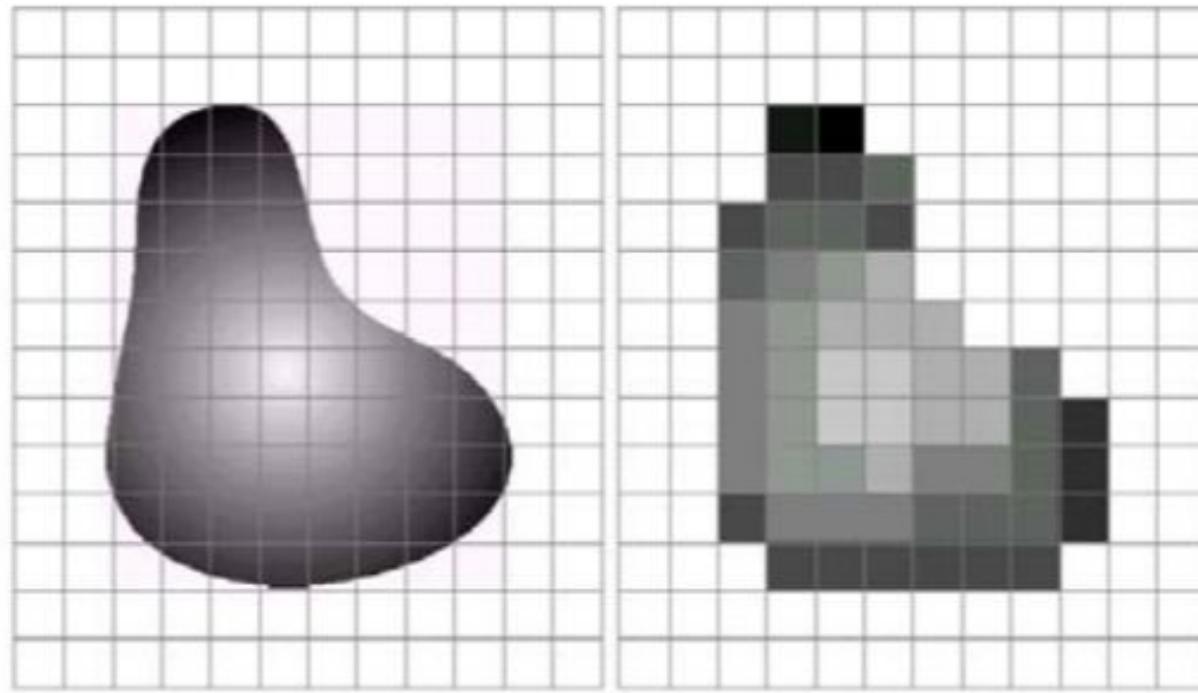


This a matrix. But from mathematical points of view we have another option

# How do machines store images?

---

We can consider image as a discrete representation of a 2D function.

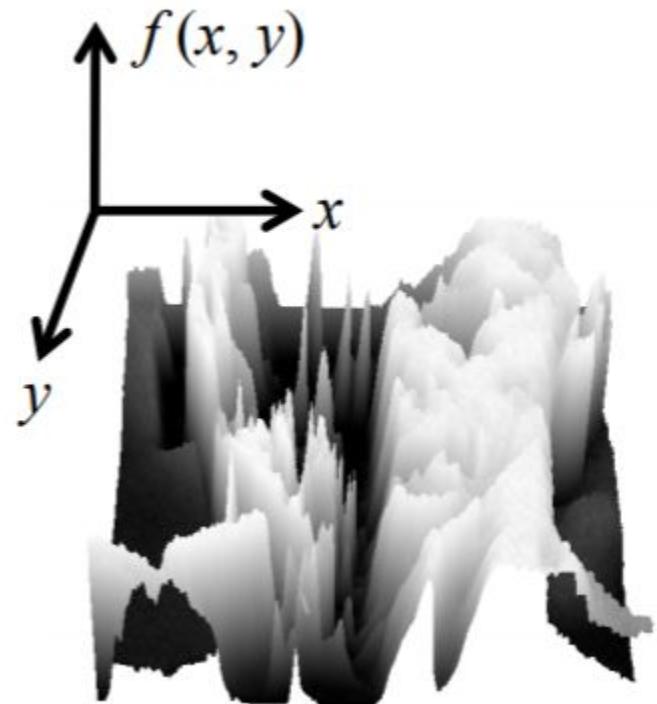


Source: D. Hoiem

# How do machines store images?

---

We can consider image as a discrete representation of a 2D function.

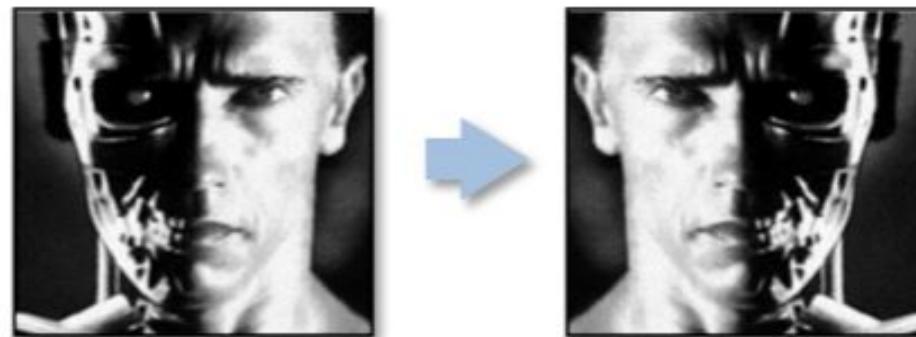
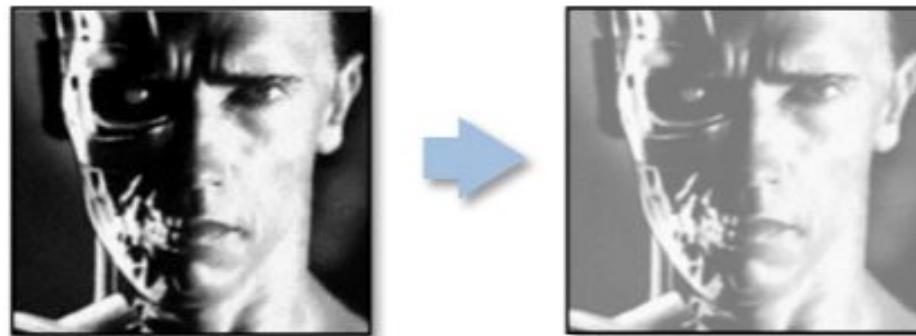


Source: N. Snavely

# How do machines store images?

---

So some standard operations can be applied

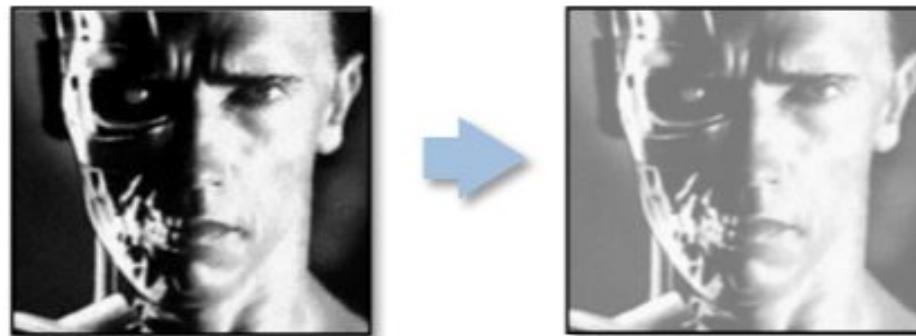


Source: N. Snavely

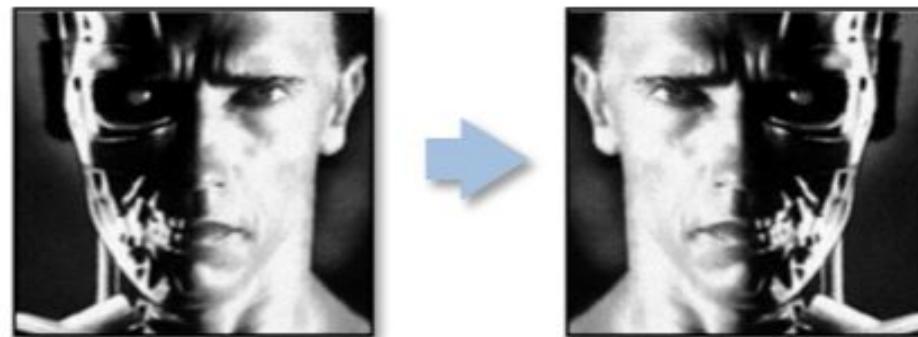
# How do machines store images?

---

So some standard operations can be applied



$$g(x,y) = f(x,y) + 20$$



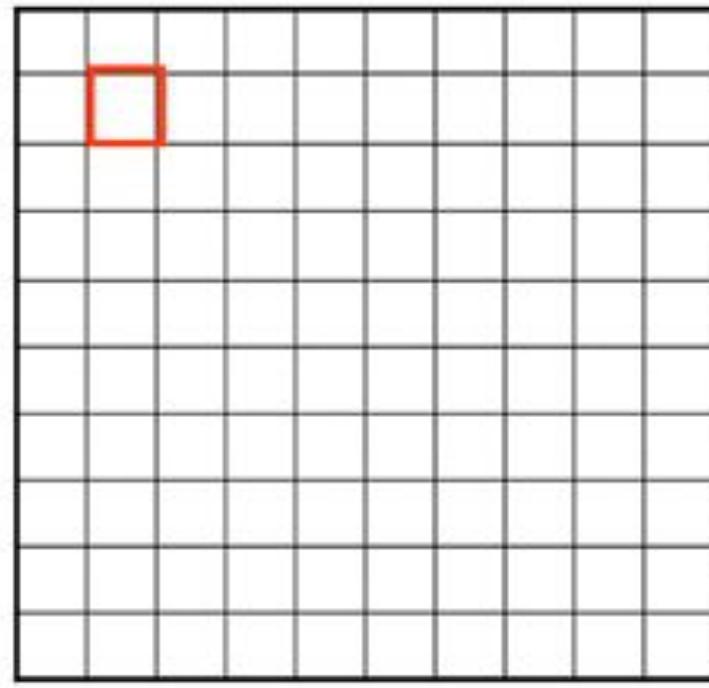
$$g(x,y) = f(-x,y)$$

Source: N. Snavely

# Convolutions: moving average example

---

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



Sources: S. Seitz, [Image filtering tutorial](#)

# Convolutions: moving average

---

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Sources: S. Seitz, [Image filtering tutorial](#)

# Convolutions: some math

---

Let us start with general 2 dimensional functions

$$(f * g)(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x', y')g(x - x', y - y')dx' dy'$$

# Convolutions: some math

---

Let us start with general 2 dimensional functions

$$(f * g)(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x', y') g(x - x', y - y') dx' dy'$$

But images are discrete representations of 2D functions

$$(f * g)(x, y) = \sum_{x'=-\infty}^{\infty} \sum_{y'=-\infty}^{\infty} f(x', y') g(x - x', y - y')$$

# Convolutions: some math

---

Let us start with general 2 dimensional functions

$$(f * g)(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x', y') g(x - x', y - y') dx' dy'$$

But images are discrete representations of 2D functions

$$(f * g)(x, y) = \sum_{x'=-\infty}^{\infty} \sum_{y'=-\infty}^{\infty} f(x', y') g(x - x', y - y')$$

And, finally, images have finite support

$$(f * g)(x, y) = \sum_{x'=0}^{X} \sum_{y'=0}^{Y} f(x', y') g(x - x', y - y')$$

# Convolutions: some math

---

Question: what is  $g$  for moving average?

$$(f * g)(x, y) = \sum_{x'=0}^X \sum_{y'=0}^Y f(x', y')g(x - x', y - y')$$

# Convolutions: some math

---

Question: what is g for moving average?

$$(f * g)(x, y) = \sum_{x'=0}^X \sum_{y'=0}^Y f(x', y')g(x - x', y - y')$$

$$\begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix}$$

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

This operation & the corresponding matrix are also called filters, kernels, convolutional matrices.

# Convolutions: moving average

---

How does it modify images?

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

# Convolutions: moving average

How does it modify images?

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



# Convolutions in 2D

---

1 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	0	0
0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	1	0
0 <small><math>\times 1</math></small>	0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature

# Convolutions in 2D: Practical exercise 1

---

Let's implement a Python function that convolves an image with a kernel.

$$(f * g)(x, y) = \sum_{x'=0}^X \sum_{y'=0}^Y f(x', y')g(x - x', y - y')$$

# Convolutions: examples of filters

---

How does it modify images?

•0	•0	•0
•0	•1	•0
•0	•0	•0

Source: J. Niebles

# Convolutions: examples of filters

How does it modify images?

•0	•0	•0
•0	•1	•0
•0	•0	•0



\*

•0	•0	•0
•0	•1	•0
•0	•0	•0

=



Source: J. Niebles

# Convolutions: examples of filters

---

How does it modify images?

•0	•0	•0
•0	•0	•1
•0	•0	•0

# Convolutions: examples of filters

How does it modify images?

•0	•0	•0
•0	•0	•1
•0	•0	•0



\*

•0	•0	•0
•0	•0	•1
•0	•0	•0

=



# Convolutions: examples of filters

---

How does it  
modify images?

$$\begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 2 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array} = ?$$

# Convolutions: examples of filters



Original

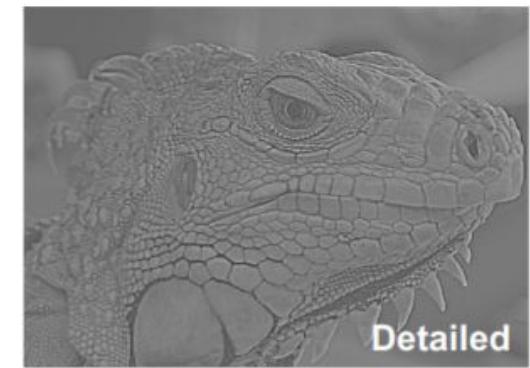
$$\begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 2 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix} - \frac{1}{9} \begin{matrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{matrix} = ?$$

(Note that filter sums to 1)

“details of the image”

$$\begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix} + \begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix} - \frac{1}{9} \begin{matrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{matrix}$$

# Convolutions: examples of filters



- Let's add it back:



# Convolutions: examples of filters



Original

$$\begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 2 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array} = ?$$

(Note that filter sums to 1)

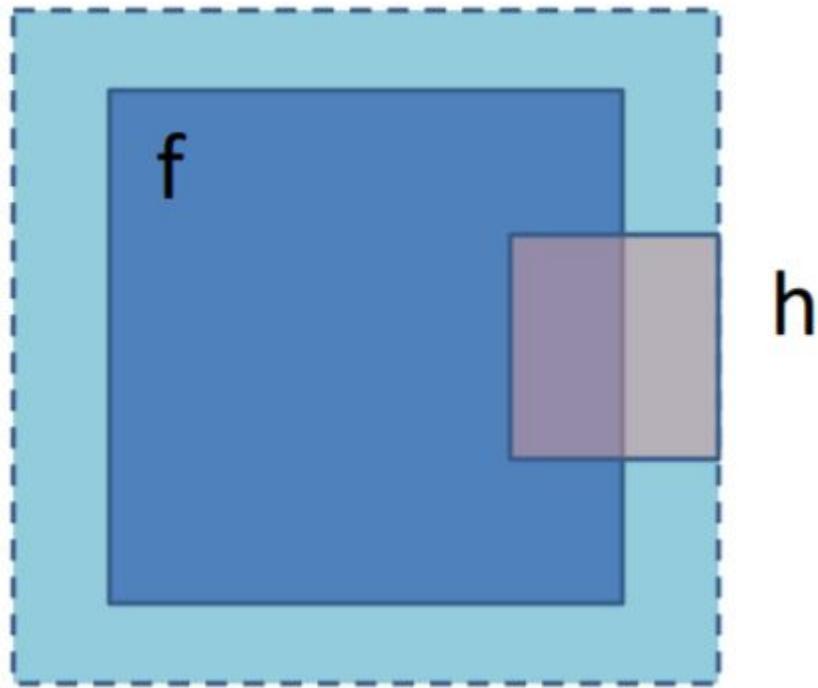


Original

$$\begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 2 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array} =$$
The result of applying the convolution operation to the original image with the first filter. The result is a grayscale image of a lizard's eye where the pupil area is slightly darker, indicating a detected feature.

# Convolutions: boundary conditions

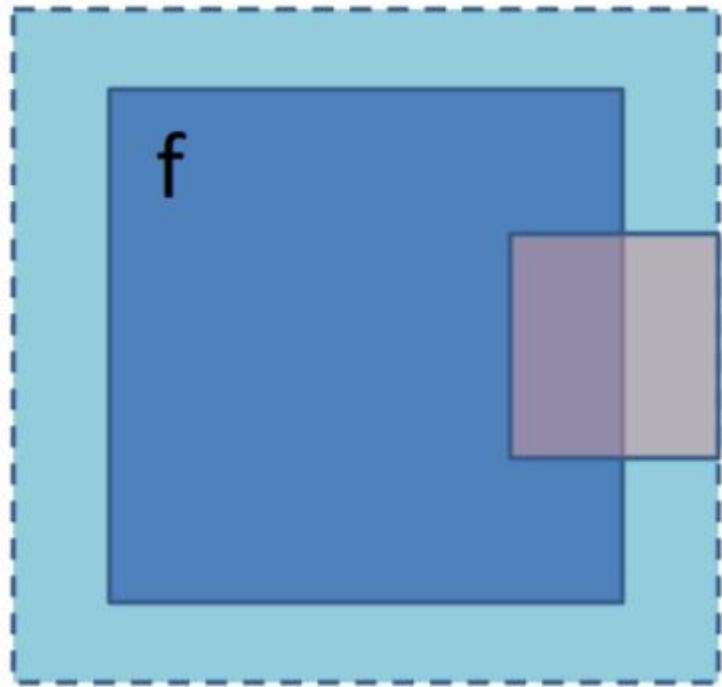
---



How to process image boundaries?

# Convolutions: boundary conditions

---



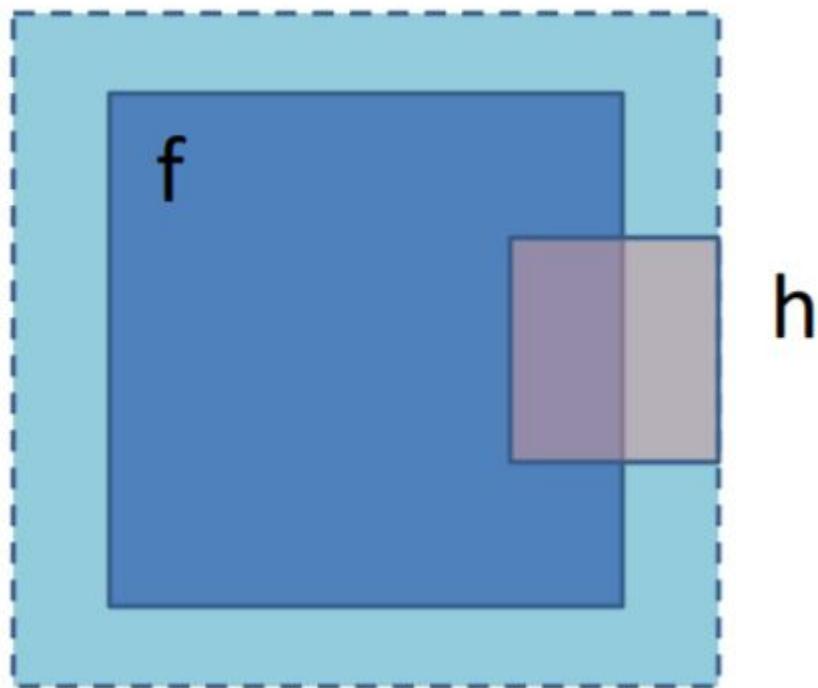
How to process image boundaries?

- **Reduce size** of the processed image
- **Add zeros** or constants
- Mirror the image

# Convolutions: boundary conditions - Task 2

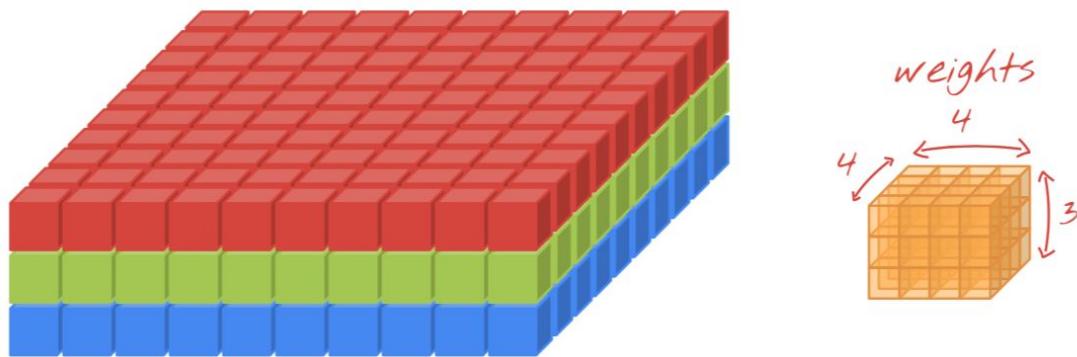
---

Let's extend our convolution function with simple zero padding.



# Convolution in 3D

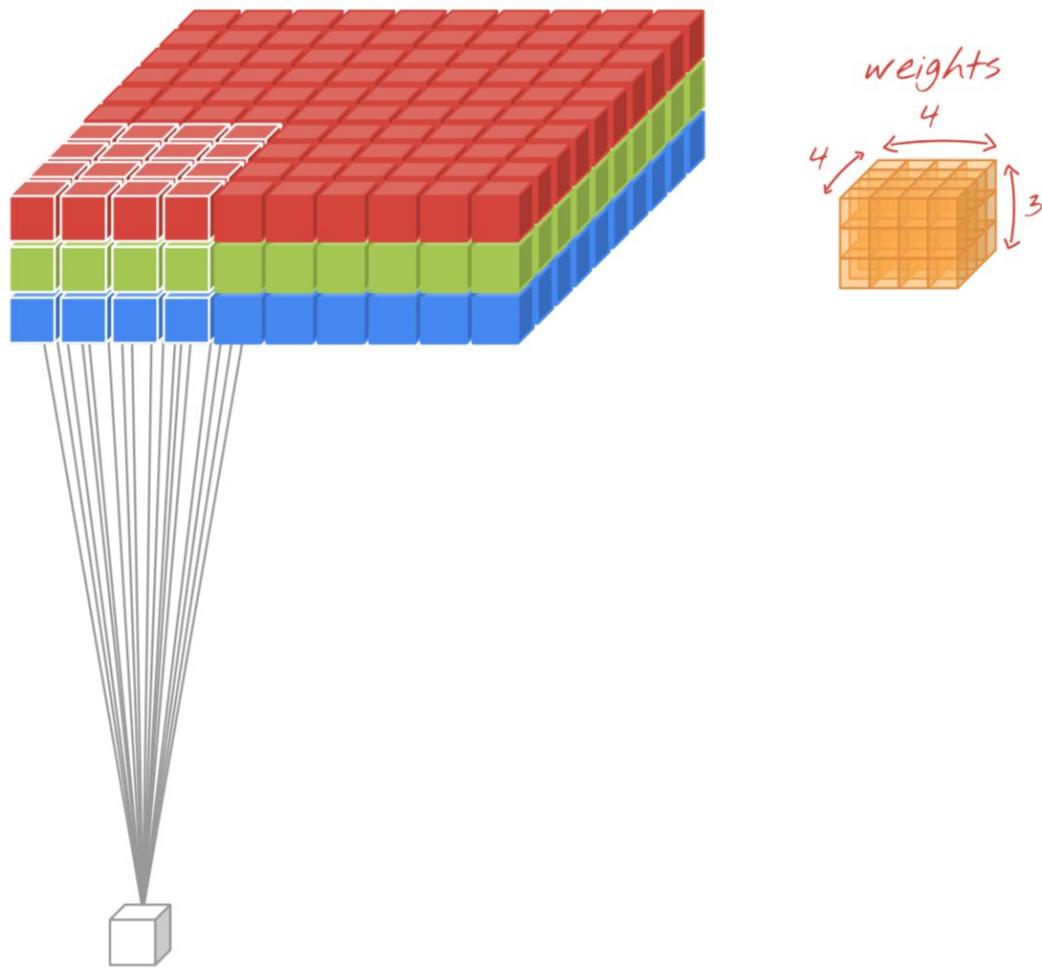
---



Source: [https://twitter.com/martin\\_gorner](https://twitter.com/martin_gorner)

# Convolution in 3D

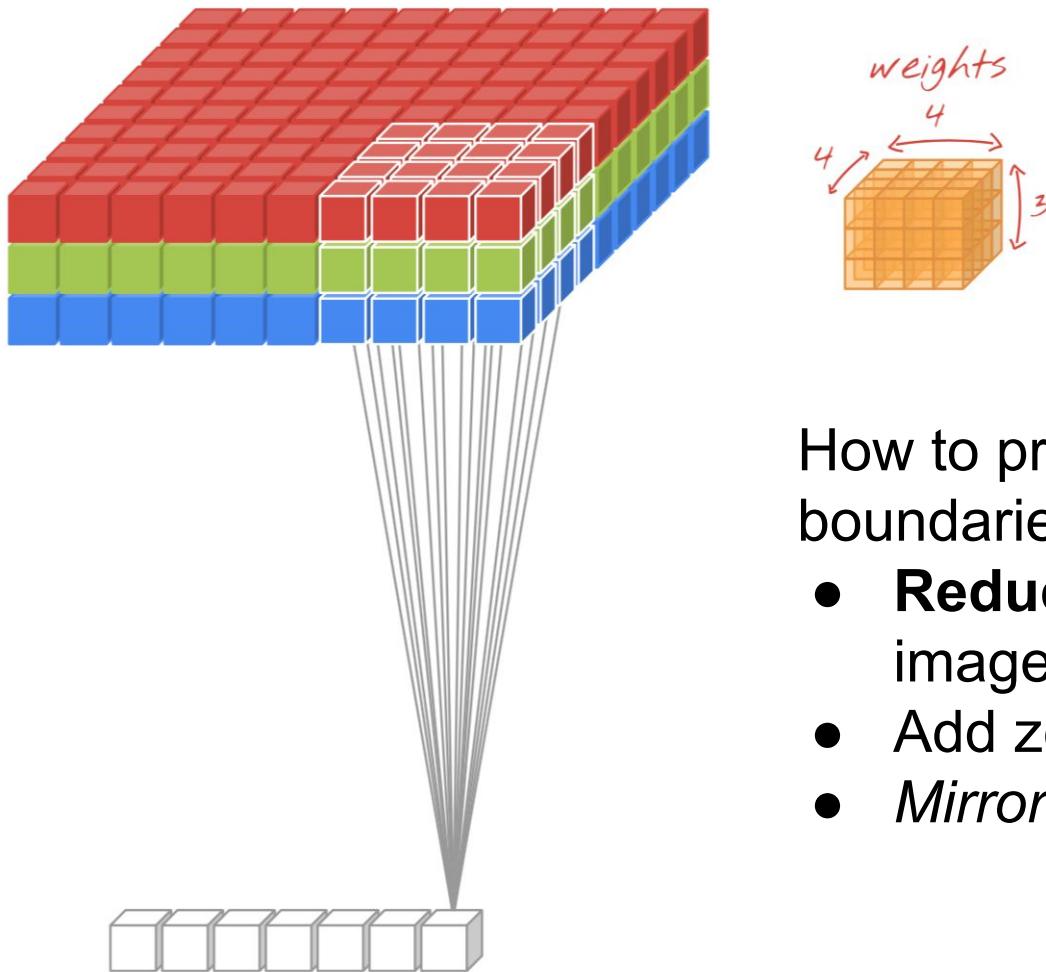
---



Source: [https://twitter.com/martin\\_gorner](https://twitter.com/martin_gorner)

# Convolution in 3D

---

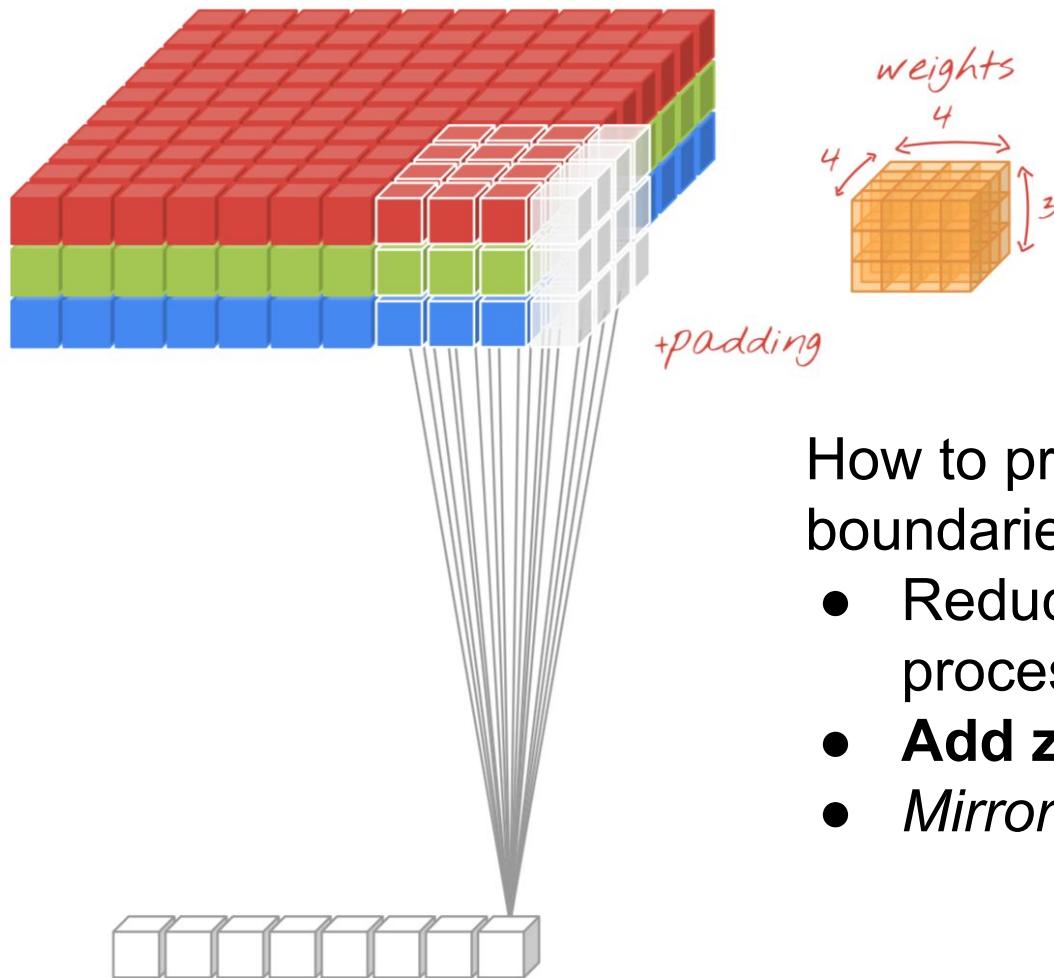


How to process image boundaries?

- **Reduce size** of the processed image
- Add zeros or constans
- *Mirror the image*

Source: [https://twitter.com/martin\\_gorner](https://twitter.com/martin_gorner)

# Convolution in 3D



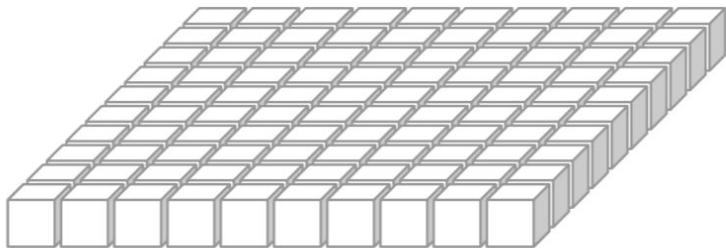
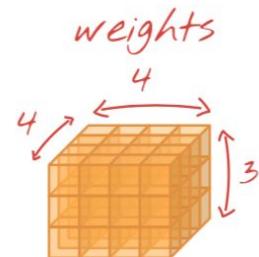
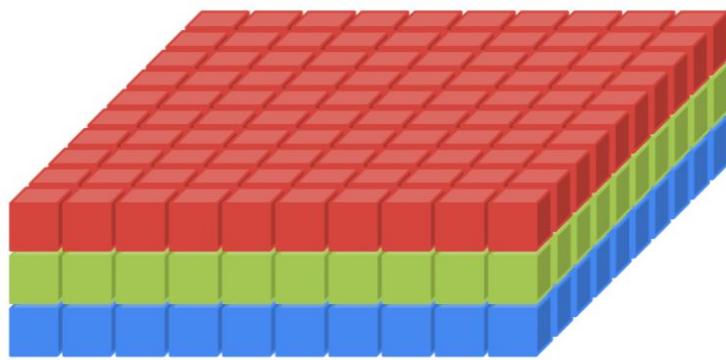
How to process image boundaries?

- Reduce size of the processed image
- **Add zeros** or constants
- *Mirror the image*

Source: [https://twitter.com/martin\\_gorner](https://twitter.com/martin_gorner)

# Convolution in 3D

---



# Outline

- Info about the course
- Computer Vision problems & applications
- Convolutions
- *Medical Computer Vision*